

Problem Set 4

*Student Name: Noah Reef***Problem 1**

Let \mathcal{H} be the set of linear hypothesis defined as

$$\mathcal{H} = \{h_a(x) := a^T x : \|a\|_2 \leq B\}$$

and suppose that for all $(x, y) \in \mathcal{D}$ we have that

$$\|x\|_2 \leq X \quad \text{and} \quad \|y\|_2 \leq Y$$

Then we can define the class of linear least square hypothesis as

$$\mathcal{L} = \{f(x) = \ell(h_a(x), y) := \frac{1}{2}(h_a(x) - y)^2 : h_a \in \mathcal{H}\}$$

Then we get that for the Radamacher complexity that

$$\begin{aligned} \mathcal{R}_S(\mathcal{H}) &= \mathbb{E}_\sigma \left[\sup_{h_a \in \mathcal{H}} \frac{1}{n} \sum_{i=1}^m \sigma_i h_a(x_i) \right] = \frac{1}{n} \mathbb{E}_\sigma \left[\sup_{\|a\|_2 \leq B} \sum_{i=1}^m \sigma_i a^T x_i \right] \\ &\leq \frac{1}{n} \mathbb{E}_\sigma \left[\sup_{\|a\|_2 \leq B} \left\| a^T \sum_{i=1}^n \sigma_i x_i \right\| \right] \\ &\leq \frac{1}{n} \mathbb{E}_\sigma \left[\sup_{\|a\|_2 \leq B} \|a\|_2 \left\| \sum_{i=1}^n \sigma_i x_i \right\|_2 \right] \\ &\leq \frac{B}{n} \mathbb{E}_\sigma \left[\left\| \sum_{i=1}^n \sigma_i x_i \right\|_2 \right] \\ &\leq \frac{B}{n} \sqrt{n} X \end{aligned}$$

and hence we have that

$$\mathcal{R}_S(\mathcal{H}) \leq \frac{BX}{\sqrt{n}}$$

lastly we can define the loss function as

$$\phi_y(z) = \ell(z, y) = \frac{1}{2}(z - y)^2$$

with $z = h_a(x)$, and compute the derivative as

$$\phi'_y(z) = (z - y)$$

and get that

$$|\phi'_y(z)| \leq |z| + |y| = |h_a(x)| + |y| \leq BX + Y$$

and hence we have that

$$\mathcal{R}_S(\mathcal{L}) \leq \frac{BX}{\sqrt{n}}(BX + Y)$$

Problem 2

```

1 import numpy as np
2 import pandas as pd
3
4 def nystrom_approx(X, sigma, r):
5     K = kernel_matrix(X, sigma)
6     A = K[0:r, 0:r]
7     B = K[r:, 0:r]
8     Z = np.block([[A], [B]])
9
10    return Z @ np.linalg.pinv(A) @ Z.T
11
12 def fourier_approx(X, sigma, r):
13     n, d = X.shape
14     # Sample random frequencies from N(0, 1/sigma^2 I)
15     W = np.random.normal(loc=0.0, scale=1.0/sigma, size=(d, r))
16     # Sample random offsets from Uniform[0, 2pi]
17     b = np.random.uniform(0, 2*np.pi, size=r)
18     # Compute the feature mapping
19     Phi = np.sqrt(2.0 / r) * np.cos(X @ W + b)
20     return Phi @ Phi.T
21
22
23
24 def kernel_matrix(X, sigma=1):
25     # Compute the squared norms of each row in X (shape: (n, 1))
26     sum_X = np.sum(X ** 2, axis=1, keepdims=True)
27
28     # Compute the pairwise squared Euclidean distances in a vectorized
29     # manner
30     dists = sum_X + sum_X.T - 2 * X @ X.T
31
32     # Compute the Gaussian kernel matrix
33     K = np.exp(-dists / (2 * sigma ** 2))
34     return K
35
36 # Create pandas DataFrame
37 df = pd.DataFrame(columns=["d", "n", "r", "Nystrom Error", "Fourier Error"])
38
39 # Test the Nystrom approximation and Fourier approximation
40 ds = [2, 4, 8, 16]
41 ns = [1024, 4096, 16384]
42 rs = [128, 512, 1024]

```

```

43 for d in ds:
44     for n in ns:
45         # Generate random data
46         mu = 0
47         sigma = 1
48         shape = (n, d)
49         X = np.random.normal(mu, sigma, shape)
50
51         # Construct total kernel matrix
52         K = kernel_matrix(X)
53
54         for r in rs:
55
56             # Construct Nystrom approximation
57             K_nys = nystrom_approx(X, 1, r)
58
59             # Construct Fourier approximation
60             K_fourier = fourier_approx(X, 1, r)
61
62             # Compute the error
63             nystrom_error = np.linalg.norm(K - K_nys) / np.linalg.norm(K)
64             fourier_error = np.linalg.norm(K - K_fourier) / np.linalg.norm
(K)
65
66
67             # Append results to DataFrame
68             df = pd.concat([df, pd.DataFrame({"d": [d], "n": [n], "r": [r
69 ], "Nystrom Error": [nystrom_error], "Fourier Error": [fourier_error]})
70 ], ignore_index=True)
71
72 # Save DataFrame to CSV
73 df.to_csv("kernel_approximation_errors.csv", index=False)
74
75 print(df.to_latex(
76     index=False,
77     float_format="%.4f",
78     column_format="|c|c|c|c|c|",
79     escape=False,
80     caption="Kernel Approximation Errors",
81     label="tab:kernel_approximation_errors"
82 ))

```

d	n	r	Nystrom Error	Fourier Error
2	1024	128	0.0033	0.1743
2	1024	512	0.0013	0.0774
2	1024	1024	0.0006	0.0639
2	4096	128	0.0025	0.1860
2	4096	512	0.0013	0.0716
2	4096	1024	0.0009	0.0812
2	16384	128	0.0023	0.1383
2	16384	512	0.0012	0.0857
2	16384	1024	0.0008	0.0693
4	1024	128	0.0698	0.3975
4	1024	512	0.0141	0.2148
4	1024	1024	0.0000	0.1406
4	4096	128	0.0785	0.3973
4	4096	512	0.0130	0.1987
4	4096	1024	0.0054	0.1463
4	16384	128	0.0684	0.4436
4	16384	512	0.0121	0.2051
4	16384	1024	0.0046	0.1589
8	1024	128	0.6445	1.6959
8	1024	512	0.3354	0.8601
8	1024	1024	0.0000	0.6053
8	4096	128	0.6695	2.0762
8	4096	512	0.3766	1.0520
8	4096	1024	0.2622	0.7291
8	16384	128	0.6157	2.1504
8	16384	512	0.3460	1.0789
8	16384	1024	0.2422	0.7586
16	1024	128	0.9349	2.8316
16	1024	512	0.7059	1.4132
16	1024	1024	0.0000	0.9979
16	4096	128	0.9839	5.6222
16	4096	512	0.9333	2.8146
16	4096	1024	0.8624	1.9895
16	16384	128	0.9953	11.0760
16	16384	512	0.9805	5.5400
16	16384	1024	0.9619	3.9167

Problem 3

Suppose we have the following Kernel regression problem $Kw = y + \eta$ where η is some noise. If K is a full-rank kernel matrix we have that

$$K = U\Sigma V^T \quad \text{with} \quad \sigma_1 \geq \sigma_2 \geq \dots \geq \sigma_n > 0$$

and hence the solution the above regression problem is given by,

$$w = K^{-1}(y + \eta) = \sum_{i=1}^n \frac{1}{\sigma_i} u_i^T (y + \eta) v_i$$

then for a low-rank approximation of our kernel matrix we have that $K_r = U_r \Sigma_r V_r^T$ and our solution vector is given by

$$w_r = \sum_{i=1}^r \frac{1}{\sigma_i} u_i^T (y + \eta) v_i = \sum_{i=1}^r \frac{1}{\sigma_i} u_i^T y v_i + \sum_{i=1}^r \frac{1}{\sigma_i} u_i^T \eta v_i = w_r^* + \delta w_r$$

then we get that

$$\|\delta w_r\| \leq \frac{\|\eta\|}{\sigma_r}$$

and so

$$\frac{\|\delta w_r\|}{\|w_r^*\|} \leq \kappa_r \frac{\|\eta\|}{\|y\|}$$

Lastly to the error is not too large we require that

$$\kappa_r \frac{\|\eta\|}{\|y\|} \leq 1 \implies \sigma_r \geq \sigma_1 \frac{\|\eta\|}{\|y\|} \implies \sigma_r \geq \sigma_1 \frac{\|\eta\|}{\|y\|}$$

Problem 4

Part a

```

1 import numpy as np
2 import pandas as pd
3 from sklearn.kernel_approximation import Nystroem
4 from sklearn.model_selection import KFold
5
6
7 def linear_regression_regularized(A, b, theta):
8     r = A.shape[1]
9     Z = A.T @ A + theta * np.eye(r)
10    y = A.T @ b
11    w = np.linalg.solve(Z, y)
12    return w
13
14 df = pd.DataFrame(columns=["Dataset", "Best Theta", "Test Error"])
15
16 for i in range(2):
17
18     # Load the in the dataset
19     fl=np.load('CSE_382M/ps_4/ps_4 code/datasets/dataset{}.npz'.format(i
20 +1))
21     xtrain = fl['xtrain'].T
22     ytrain = fl['ytrain'].T

```

```
23 # regularization parameter theta
24 thetas = [1e-4,1e-3,1e-2,1e-1,1,10,100]
25
26 # errors
27 err = {}
28
29 # Create K-Fold Cross-Validation split of the training dataset
30 kf = KFold(n_splits=5, shuffle=True, random_state=42)
31 kf.get_n_splits(xtrain)
32 for train_index, test_index in kf.split(xtrain):
33     x_train, x_test = xtrain[train_index], xtrain[test_index]
34     y_train, y_test = ytrain[train_index], ytrain[test_index]
35
36     # Loop over the regularization parameters
37     for theta in thetas:
38         feature_map = Nystroem(kernel='rbf', gamma=1, n_components
=100)
39         x_train_transformed = feature_map.fit_transform(x_train)
40         x_test_transformed = feature_map.transform(x_test)
41
42         # Train the model
43         w = linear_regression_regularized(x_train_transformed, y_train
, theta)
44
45         # Make predictions
46         y_pred = x_test_transformed @ w
47         # Compute the error
48         error = np.mean((y_test - y_pred) ** 2)
49         # Store the error
50         if theta not in err:
51             err[theta] = []
52         err[theta].append(error)
53     # Compute the average error for each theta
54     for theta in err:
55         err[theta] = np.mean(err[theta])
56
57     # pick best theta
58     best_theta = min(err, key=err.get)
59
60     # test the best theta
61     yTest = fl['ytest'].T
62     xTest = fl['xtest'].T
63     xtrain = fl['xtrain'].T
64     ytrain = fl['ytrain'].T
65
66     feature_map = Nystroem(kernel='rbf', gamma=1, n_components=100)
67     xtrain_transformed = feature_map.fit_transform(xtrain)
68     xTest_transformed = feature_map.transform(xTest)
69     w = linear_regression_regularized(xtrain_transformed, ytrain,
best_theta)
70     y_pred = xTest_transformed @ w
71
72     # Compute the error
73     error = np.mean((yTest - y_pred) ** 2)
```

```
74     # Append the results to the DataFrame
75     df = pd.concat([df, pd.DataFrame({"Dataset": [i+1], "Best Theta": [
best_theta], "Test Error": [error]})], ignore_index=True)
76
77 df.to_csv("dataset_errors.csv", index=False)
78
79 print(df.to_latex(index=False,
80                  float_format="{:.3E}".format,
81 ))
```

Part b

Dataset	Best Theta	Test Error
1	1.000E-03	9.973E-09
2	1.000E+00	1.226E-09